

# Project 2

(MCP: *Ghost-in-the-shell*)

CIS 415 - Operating systems

2021 Spring- Prof. Allen Malony

Due date: **11:59 pm, Sunday, May 16th, 2021**

## Introduction:

The British philosopher, Gilbert Ryle, in *The Concept of Mind* criticizes the classical Cartesian rationalism that the mind is distinct from the body (known as mind-body duality). He argues against the “dogma of the ghost in the machine” as an independent non-material entity, temporarily inhabiting and governing the body. In *The Ghost in the Machine*, Arthur Koestler furthers the position that the mind-body duality is a whole and not a part. The manga series *Ghost in the Shell* by Masamune Shirow is an homage to Koestler, explores the nature of the mind, and whether it can exist without the body. When I was going to UCLA for my master's degree, I was an intern at Burroughs Corporation. Sadly, Burroughs no longer exists, having been acquired by Hewlett-Packard. It was a major player in computing systems, and my job was to write the *MCP*:

<http://www.burroughs.com/mcp>

<http://www.burroughs.com/mcp>

The MCP might be considered one of the most innovative in mainframe operating systems, but it was innovative in many ways. It was the first OS to support multiple processors, the first OS to support a large number of users, and the first OS written exclusively in a high-level language. MCP executed “jobs” with each containing one or more tasks. Tasks within a job could run sequentially or in parallel. Logic could be implemented at the job level to control the flow of a job by writing in the MCP’s workflow control language (WFL). Once all tasks in a job completed, the job itself was done. In some respects, we might regard these features of MCP as now being provided through shell scripts that access an operating system’s services. I often wonder whether the ghosts of the modules I wrote for MCP continue to live in the shells being used today.

## Project Details:

In this project, you will implement the *MCP Ghost in the Shell* (MCP for short) whose primary

job is to launch a pool of sub-processes that will execute commands given in an input file. The MCP will read a list of commands (with arguments) to run from a file, it will then start up and run the process that will execute the command, and then schedule the processes to run concurrently in a time-sliced manner. In addition, the MCP will monitor the processes, keeping track of how the processes are using system resources. There are a total of 4 parts to the project, each building on the other. The objective is to give you a good introduction to processes, signals, signal handling, and scheduling. Each part is a complete program by itself and **must** be saved in a separate \*.c file in different directories with their own make file.

**Note: For part 1, Do not test to spawn more than 7 processes.**

## Part 1: MCP Launches the Workload

For part 1 of this project, you will be developing the first version of the MCP such that it can launch the workload and

### Program Requirements

Your MCP v1.0 must

1. Read the program arguments from the command line and work just like the file mode from the command line and its arguments.
2. For each command in the input file, run the command using some variable.
  - a. **fork(2)**
  - b. **One of the following:**
    - i. **wait(2)**
3. Once all of the processes have been launched, the process to terminate using one of the following:
  - a. **wait(2)**
4. After all processes have been launched, the process to terminate using one of the following:
  - a. **exit()**


Fig. 1 presents the pseudocode for launching processes:

```

for (int i = 0; i < line_number; i++)
{
    pid_array[i] = fork();
    if (pid_array[i] < 0)
    {
        //error handling
    }

    if (pid_array[i] == 0)
    {
        if (execvp (path, arg) == -1)
        {
            error handling
        }
    }
}

```



### Remarks:

To make things simple, you should use the MCP (i.e. the VM) to run your program. This is because many things can go wrong. You should also be aware of these system calls. Also, you will need to get the arguments from the “workload” file and get the arguments. The arguments can be anything that will run in the same environment in which the MCP is launched. A set of commands should be used to launch the MCP. You should also be aware of the influence on your final grade so take some time to think about how your program represents what is going on at this stage. What is expected is that your output will be somewhat jumbled. This can be an excellent indicator that your code is working as intended. If you find that your messages are all synchronized then you know that something is going wrong.

## Part 2: MCP Controls the Workload

Successful completion of Part 1 will give you a basic working MCP. However, if we just wanted to run all the workload programs at the same time, we might as well use a shell script. Rather, our ultimate goal is to schedule the programs in the workload to run in a time-shared manner. Part 2 will take the first step to allow the MCP to gain control for this purpose. This will be completed in two core steps:

### Step 1:

Step 1 of Part 2 will implement a way for the MCP to stop all forked (MCC) child processes right before they call `exec()`. This gives the MCP back control before any of the workload programs are launched. The MCP will then wait for a `SIGUSR1` signal before launching the workload programs. The `sigwait()` system call will be used to wait for the `SIGUSR1` signal. The `sigwait()` system call is *blocking*, meaning that the process will wait for the signal. *As the `exec()` call, it is non-blocking*, meaning that the process will not wait for the signal. When the MCP is in a state where each child process is selected to be launched, the MCP will send the `SIGUSR1` signal to it, which will launch the associated workload program.

### Step 2:

Step 2 will immediately stop (using the `SIGSTOP` signal). This is the mechanism that sends a `SIGSTOP` signal and typing Ctrl-Z to stop a running process to bring it back to the command line. The `SIGCONT` signal is like

### Program Requirements:

Thus, in Part 2, you will take your MCP v1.0 and implement both Step 1 and 2 to create MCP v2.0. Your MCP v2.0 **must** do the following:

1. Immediately after each program is created using the **fork()** system call, the forked MCP child process waits until it receives a **SIGUSR1** signal before calling **exec()**.
2. After all of the MCP child processes have been forked and are now waiting, the MCP parent process must send each child process a **SIGUSR1** signal (at the same time) to wake them up. Each child process will then return from the **sigwait()** and call **exec()** to run the workload program.
3. After all of the workload programs have been launched and are now executing, the MCP



must send each child process a `SIGSTOP` signal to suspend them.

4. After all of the child processes have been suspended, the MCP must send each child process a `SIGCONT` signal to wake them up.
5. Again, once all of the processes are back up and running, the MCP must wait for each child process to terminate. After all processes have terminated, the MCP will exit and free any allocated memory.

### Remarks:

MCP 2.0 demonstrates that we can control the suspending and continuing of processes. You should test out that things are working properly. One way to do this is to create messages that indicate what steps the MCP is presently taking and for which child process. The console output of your program has a heavy influence on your final grade so take some time to think about how your program should look.

Again, a set of example messages is provided, but you should also construct your own. For more information, see the section on debugging. We recommend that you read the man pages for the system calls used in the example program.

### Part 3: MCP Scheduler

Now that the MCP is working, you will implement a scheduler. The simplest policy is to equally share the CPU among all child processes (e.g., 1 second). The scheduler will run for a certain amount of time (e.g., 1 second). After its time slice has completed, the MCP will suspend the process. The MCP decides which is the next process to run.

### Program Requirements

MCP 2.0 knows how to run a process. For Part 3 of this project, you will be upgrading your MCP v2.0 to include process scheduling. *Note: if a child process is running, it is still the case that the MCP is running “concurrently” with it.* Thus, one way to approach the problem of process scheduling is for the MCP to poll the system time to determine when the time slice is expended. This is inefficient and not as accurate as it can be.

Thus, we will be using signal processing to implement a more intelligent way of process scheduling. In general, this is done by setting an alarm using the `alarm(2)` system call. This tells the operating system to deliver a `SIGALRM` signal after some specified time. In general, signal handling is done by registering a signal handling function with the operating system. When the signal is delivered, the MCP will be interrupted and the signal handling function will be executed. The following must happen after the alarm signal is received by the MCP:

1. The MCP will suspend the currently running workload process using `SIGSTOP`.

2. The MCP decides on the next workload process to run, and sends it a `SIGCONT` signal.
3. The MCP will reset the alarm, and continue with whatever else it was doing.

### Remarks:

Your new and improved MCP v3.0 is now a working workload process scheduler. However, you need to take care of a few things. For instance, there is the question of how to determine if a workload process is still executing. At some point (we hope), the workload process is going to terminate. Remember, this workload process is a child process of the MCP. How does the MCP know that the workload process has terminated? In MCP v2.0, we just called `wait()`. Is that sufficient now? Be careful. (**Note:** You cannot use the return value of `waitpid()` to determine if a process has terminated. There are multiple correct answers to this issue. So give it some thought and choose appropriately. Again, be careful.)

Finally, please demonstrate your program and provide some feedback. Having a demonstration is one way to do this. The console output should be at least a B grade so take some time to think about how to present it.

### Part 4: MCP Knowledge

With MCP v3.0, the workload processes get an “equal” share of the processor. Each workload program it reads in. You should also write your own simple test program to verify that the scheduler is proceeding by looking in the `/proc` directory.

### Program Requirements

In Part 4, you will write a program that conveys some information about the processes consuming. You will pick out a minimum of 5 properties for each process and display them on your terminal. You will also update these informations according to the time slice you assigned from part 3. This may include information about execution time, memory used, and I/O. It is up to you to decide what to look at, analyze, and present. Do not just dump out everything in `/proc` for each process. The objective is to give you some experience with reading, interpreting, and analyzing process information. Your MCP 4.0 must output the analyzed process information for every child process each time the scheduler completes a cycle. Of course, the format of the output is important too. One thought is to do something similar to what the Linux `top(1)` program does (i.e. format the information into a table that constantly updates each cycle).

### Project Remarks:

### Error Handling:

All system call functions that you use will report errors via the return value. As a general rule, if the return value is less than zero, then an error has occurred and `errno` is set accordingly. You must check your error conditions and report errors. To expedite the error checking process, we will allow you to use the `perror(3)` library function. Although you are allowed to use `perror`, it does not mean that you should report errors with voluminous verbosity. Report fully but concisely.

## Memory Errors:

You are required to check the most important one. Code that is not deduced. Fortunately, `valgrind`, while quite useful, is not perfect. It can't find those that are located by variables in your code: especially



task, but a very  
I have marks  
Remember that  
find and fix any bugs  
y memory error in

## Developing Your Co

The best way to develop a strategy is the one provided to you. This gives you the benefit of taking a risky or hazardous, so-called state.

chine image  
restart. This also gives  
something potentially  
oll back to a safe

## Part 5: Extra Credit

When the MCP schedules a workload process to run, it assumes that the process will actually execute the entire time slice. However, suppose that the process is doing I/O, for example, waiting for the user to enter something on the keyboard. In general, workload processes could be doing different things and thus have different execution behaviors. Some processes may be compute bound while others may be I/O bound. If the MCP knew something about process behavior, it is possible that the time slice could be adjusted for each type of process. For instance, I/O bound processes might be given a little less time and compute bound processes a bit more. By adjusting the time slice, it is possible that the entire workload could run more efficiently. Part 5 is to implement some form of adjustable scheduler that uses process information to set process specific time intervals. These time intervals should then be shown in a new column in your tabular display along with another column indicating the type of work detected.

*Note: Since this is for extra credit, we reserve the right to be strict here. The extra points may come in an all-or-none fashion. Also be careful not to introduce new errors or bugs into your program.*

■





For a project to be accepted, the project must contain the following 5 files and meet the following requirements: (The naming conventions listed below **must** be followed. Additionally you must use the C programming language for this assignment. No projects written in another programming language will be accepted.)

**part2.c:** This is the c-file that contains your MCP v2.0. It must meet the specifications described in Part 2 of the project details.

**part4.c:** This is the c file that contains the functions and configurations described in Part 4 of the project

**MCP.h (optional):** You can add whatever you wish to MCP.h and structs. You can completely.

**Makefile:** Your project files should be named with the following names: **part1**, **part2**, **part3**, **part4**, **part5**, **part6**, **part7**, **part8**, **part9**, **part10**, **part11**, **part12**, **part13**, **part14**, **part15**, **part16**, **part17**, **part18**, **part19**, **part20**. Do not include the names of the files in the Makefile. As

**Note:** Additionally, you are allowed to add any other \*.h and \*.c files you wish. However, when we run your code we will only be running the part\* files. Make sure your code runs in the VM before submission.

Once your project is done, do the following:

1. Open a terminal and navigate to the project folder. Compile your code in the VM with the -g flag.
2. Run your code and take screenshots of the output as necessary (of each part).
3. Create valgrind logs of each respective part:
  - a. **“valgrind --leak-check=full --tool=memcheck ./a.out > log\*.txt 2>&1 ”**
4. Create a Tar containing only the files/logs/screenshots and submit it onto Canvas.
5. Submit a .pdf of your project report separately onto canvas. (i.e. you should upload two things: your report and the tar.gz)

Valgrind can help you spot memory leaks in your code. As a general rule any time you allocate memory you must free it. Points will be deducted in both the labs and the project for memory leaks so it is important that you learn how to use and read Valgrind's output. See

(<https://valgrind.org/>) for more details. Finally, please abide by the submission guidelines, we

will be grading your p result in lost points.



## Grading Rubric:

Part	Points	Description
Part 1	15	5 error handling 5 wait for processes to finish 10 correct processes spawned
Part 2	30	10 correct implementation of sigwait() 10 correct sequence of signal sent 5 sleep() called in between signals for understandable format 5
Part 3	30	10 10 20 10 25
Part 4	15	4 pr 15 -5
Valgrind	5	1
Report	5	R (pages)
Part 5	10	C

**Note:** some sections may have more detail points than the total points, meaning there are more than 1 way you can get a 0 in that section.

1. 0 if your program does not compile.
2. 0 if your program does not work in the required environment.
3. 5 points deduction if your makefile does not work.

### Late Homework Policy:

- 10% penalty (1 day late)
- 20% penalty (2 days late)
- 30% penalty (3 days late)

- 100% penalty (>3 days late) (i.e. no points will be given to homework received after 3 days)

